

Breast Cancer prediction using Neural Networks

Ilija Doknić
Stefan Komarica







Dataset



- Binary classification prediction for determining is breast cancer malignant or benign type
- Data taken from <https://www.kaggle.com/datasets/yasserh/breast-cancer-dataset>





Initial analysis

diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	...
M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	...
M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	...
M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	...
M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	...
M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	...











Neural network





Single neuron



$$z = w^T x + b$$

$$a = \sigma(z)$$


```
class Dense:
    ...

    def forward(self, inputs):
        self.inputs = inputs
        self.output = np.dot(inputs, self.weights) + self.bias

class ReLU:
    def forward(self, inputs):
        self.inputs = inputs
        self.output = np.maximum(0, inputs)

class Sigmoid:
    def forward(self, inputs):
        self.output = 1 / (1 + np.exp(-np.array(inputs)))
```



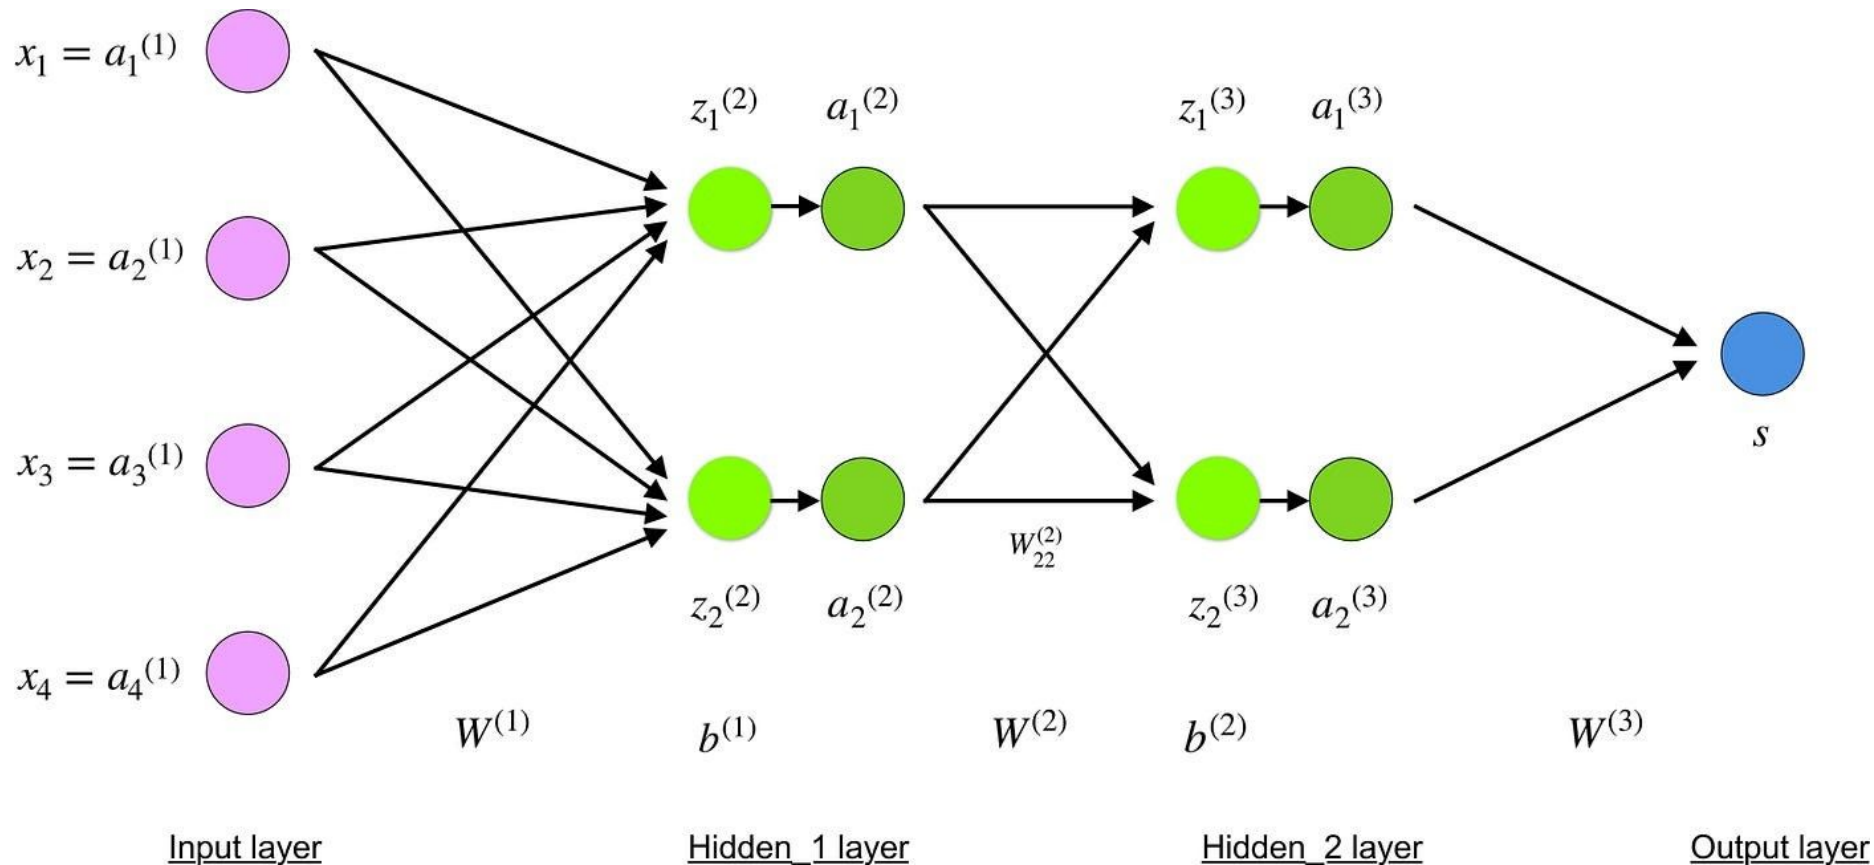
LogLoss function

$$\text{Log Loss}(y, p) = -\frac{1}{N} \sum_{i=1}^N (y_i \log(p_i) + (1 - y_i) \log(1 - p_i))$$




Implementation

```
class LogLoss(Loss):  
    def forward(self, y_pred, y_true):  
        y_pred = np.clip(y_pred, 1e-7, 1 - 1e-7)  
        return (-1/len(y_true))*\  
            (np.dot(y_true.T, np.log(y_pred))+\  
             np.dot((1-y_true).T, np.log(1-y_pred)))
```



Back propagation






Implementation


```
class Dense:
    ...
    def backward(self, deriv_in):
        self.deriv_weights = np.dot(self.inputs.T, deriv_in)
        self.deriv_bias = np.sum(deriv_in, axis=0, keepdims=True)
        self.deriv_out = np.dot(deriv_in, self.weights.T)

class Sigmoid:
    ...
    def backward(self, deriv_in):
        self.deriv_out = self.output*(1-self.output)*deriv_in
```

SGD without momentum



SGD with momentum



Adam optimization algorithm

```
class Adam():
    def update_params(self, layer):

        if not hasattr(layer, 'weight_cache'):
            layer.weight_cache = np.zeros_like(layer.weights)
            layer.weight_momentum = np.zeros_like(layer.weights)
            layer.bias_cache = np.zeros_like(layer.bias)
            layer.bias_momentum = np.zeros_like(layer.bias)
        # Momentum part of Adam
        layer.weight_momentum = self.beta1*layer.weight_momentum + (1-self.beta1)*layer.deriv_weights
        layer.bias_momentum = self.beta1*layer.bias_momentum + (1-self.beta1)*layer.deriv_bias

        weight_momentum_corrected = layer.weight_momentum/(1 - self.beta1**(self.iteration+1))
        bias_momentum_corrected = layer.bias_momentum/(1 - self.beta1**(self.iteration+1))

        layer.weight_cache = self.beta2*layer.weight_cache + (1-self.beta2)*layer.deriv_weights**2
        layer.bias_cache = self.beta2*layer.bias_cache + (1-self.beta2)*layer.deriv_bias**2

        weight_cache_corrected = layer.weight_cache/(1-self.beta2**(self.iteration+1))
        bias_cache_corrected = layer.bias_cache/(1-self.beta2**(self.iteration+1))

        layer.weights -= self.current_learning_rate*weight_momentum_corrected/(np.sqrt(weight_cache_corrected+self.epsilon))
        layer.bias -= self.current_learning_rate*bias_momentum_corrected/(np.sqrt(bias_cache_corrected+self.epsilon))

    def post_update_params(self):
        self.iteration += 1
```


Repeat multiple times

```
for epoch in range(21):
    dense_1.forward(X_test)
    relu_1.forward(dense_1.output)
    dense_2.forward(relu_1.output)
    relu_2.forward(dense_2.output)
    dense_3.forward(relu_2.output)
    sigmoid.forward(dense_3.output)
    current_loss = loss.forward(sigmoid.output, y_test)
    y_pred_test = np.array(sigmoid.output > 0.5, dtype = int)

    dense_1.forward(X_train)
    relu_1.forward(dense_1.output)
    dense_2.forward(relu_1.output)
    relu_2.forward(dense_2.output)
    dense_3.forward(relu_2.output)
    sigmoid.forward(dense_3.output)
    current_loss = loss.forward(sigmoid.output, y_train)
    y_pred = np.array(sigmoid.output > 0.5, dtype = int)

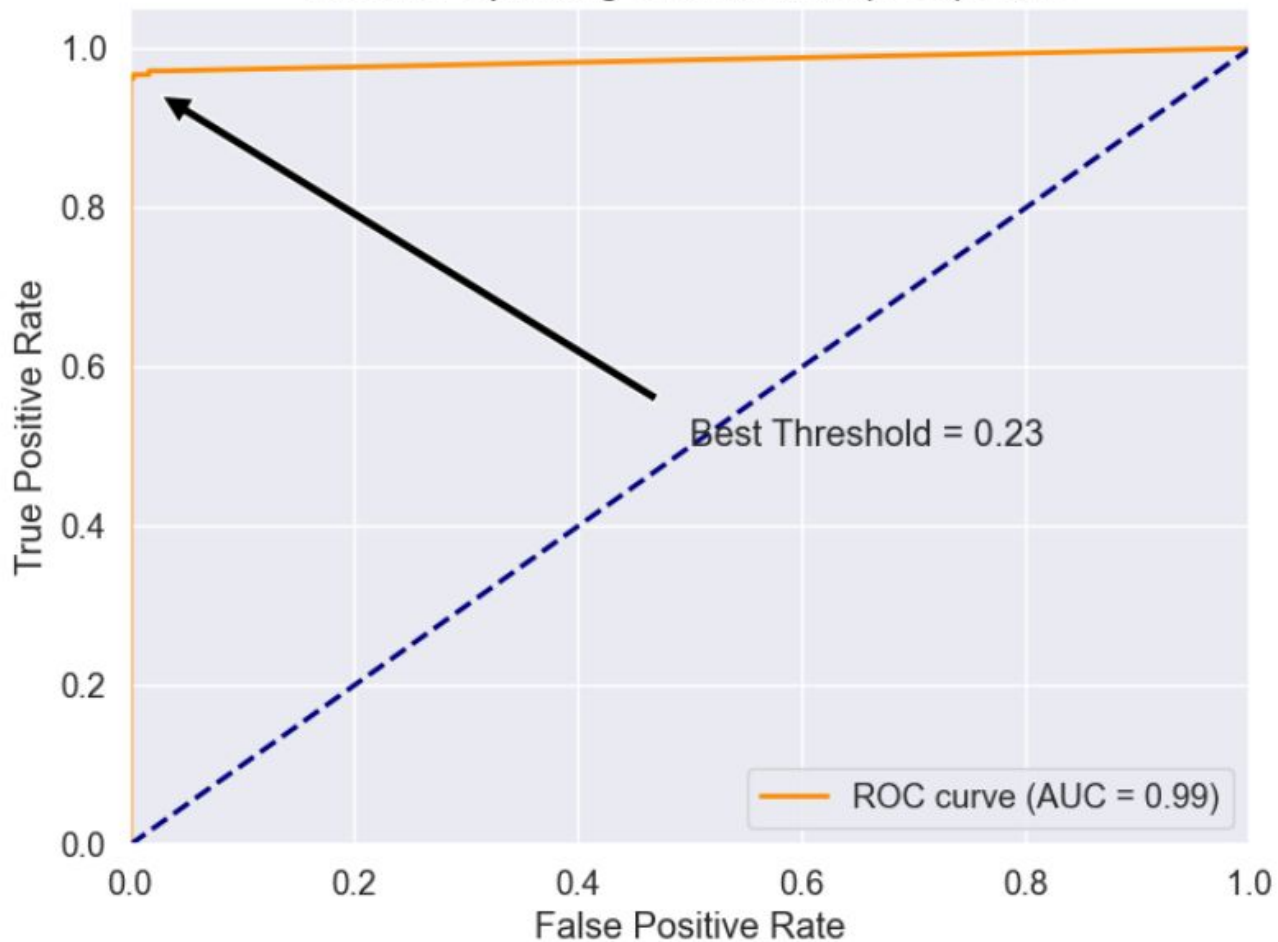
    loss.backward(sigmoid.output, y_train)
    sigmoid.backward(loss.deriv_out)
    dense_3.backward(sigmoid.deriv_out)
    relu_2.backward(dense_3.deriv_out)
    dense_2.backward(relu_2.deriv_out)
    relu_1.backward(dense_2.deriv_out)
    dense_1.backward(dense_2.deriv_out)

    optimizer.update_params(dense_3)
    optimizer.update_params(dense_2)
    optimizer.update_params(dense_1)
```

Accuracy and loss results

```
Epoch: 1, Accuracy: 0.8451, Val-Accuracy: 0.8392, Loss: 0.6931
Epoch: 2, Accuracy: 0.6268, Val-Accuracy: 0.6294, Loss: 0.6606
Epoch: 3, Accuracy: 0.6268, Val-Accuracy: 0.6294, Loss: 0.6603
Epoch: 4, Accuracy: 0.6268, Val-Accuracy: 0.6294, Loss: 0.6598
Epoch: 5, Accuracy: 0.6268, Val-Accuracy: 0.6294, Loss: 0.6583
Epoch: 6, Accuracy: 0.6268, Val-Accuracy: 0.6294, Loss: 0.6507
Epoch: 7, Accuracy: 0.6948, Val-Accuracy: 0.6853, Loss: 0.5904
Epoch: 8, Accuracy: 0.9366, Val-Accuracy: 0.9371, Loss: 0.3593
Epoch: 9, Accuracy: 0.9413, Val-Accuracy: 0.951, Loss: 0.2293
Epoch: 10, Accuracy: 0.9742, Val-Accuracy: 0.972, Loss: 0.1778
Epoch: 11, Accuracy: 0.9554, Val-Accuracy: 0.958, Loss: 0.1539
Epoch: 12, Accuracy: 0.9718, Val-Accuracy: 0.979, Loss: 0.1253
Epoch: 13, Accuracy: 0.9765, Val-Accuracy: 0.979, Loss: 0.1161
Epoch: 14, Accuracy: 0.9765, Val-Accuracy: 0.979, Loss: 0.108
Epoch: 15, Accuracy: 0.9742, Val-Accuracy: 0.979, Loss: 0.1009
Epoch: 16, Accuracy: 0.9789, Val-Accuracy: 0.986, Loss: 0.095
Epoch: 17, Accuracy: 0.9789, Val-Accuracy: 0.986, Loss: 0.0867
Epoch: 18, Accuracy: 0.9836, Val-Accuracy: 0.986, Loss: 0.0813
Epoch: 19, Accuracy: 0.9836, Val-Accuracy: 0.986, Loss: 0.078
Epoch: 20, Accuracy: 0.9812, Val-Accuracy: 0.986, Loss: 0.0755
Epoch: 21, Accuracy: 0.9836, Val-Accuracy: 0.986, Loss: 0.0736
```

Receiver Operating Characteristic (ROC) Curve



Confusion Matrix

